

C declaration	Intel data type	Assembly-code suffix	Size (bytes)
char	Byte	b	1
short	Word	w	2
int	Double word	l	4
long	Quad word	q	8
char *	Quad word	q	8
float	Single precision	s	4
double	Double precision	l	8

63	31	15	7	0	
%rax	%eax	%ax	%al		Return value
%rbx	%ebx	%bx	%bl		Callee saved
%rcx	%ecx	%cx	%cl		4th argument
%rdx	%edx	%dx	%dl		3rd argument
%rsi	%esi	%si	%sil		2nd argument
%rdi	%edi	%di	%dil		1st argument
%rbp	%ebp	%bp	%bpl		Callee saved
%rsp	%esp	%sp	%spl		Stack pointer
%r8	%r8d	%r8w	%r8b		5th argument
%r9	%r9d	%r9w	%r9b		6th argument
%r10	%r10d	%r10w	%r10b		Caller saved
%r11	%r11d	%r11w	%r11b		Caller saved
%r12	%r12d	%r12w	%r12b		Callee saved
%r13	%r13d	%r13w	%r13b		Callee saved
%r14	%r14d	%r14w	%r14b		Callee saved
%r15	%r15d	%r15w	%r15b		Callee saved

Type	Form	Operand value	Name
Immediate	$\$Imm$	$Imm$	Immediate
Register	$r_a$	$R[r_a]$	Register
Memory	$Imm$	$M[Imm]$	Absolute
Memory	$(r_a)$	$M[R[r_a]]$	Indirect
Memory	$Imm(r_b)$	$M[Imm + R[r_b]]$	Base + displacement
Memory	$(r_b, r_i)$	$M[R[r_b] + R[r_i]]$	Indexed
Memory	$Imm(r_b, r_i)$	$M[Imm + R[r_b] + R[r_i]]$	Indexed
Memory	$(, r_i, s)$	$M[R[r_i] \cdot s]$	Scaled indexed
Memory	$Imm(, r_i, s)$	$M[Imm + R[r_i] \cdot s]$	Scaled indexed
Memory	$(r_b, r_i, s)$	$M[R[r_b] + R[r_i] \cdot s]$	Scaled indexed
Memory	$Imm(r_b, r_i, s)$	$M[Imm + R[r_b] + R[r_i] \cdot s]$	Scaled indexed

1	movl \$0x4050,%eax	Immediate--Register, 4 bytes
2	movw %bp,%sp	Register--Register, 2 bytes
3	movb (%rdi,%rcx),%al	Memory--Register, 1 byte
4	movb \$-17, (%esp)	Immediate--Memory, 1 byte
5	movq %rax, -12(%rbp)	Register--Memory, 8 bytes

Instruction	Effect	Description
leaq $S, D$	$D \leftarrow \&S$	Load effective address
inc $D$	$D \leftarrow D+1$	Increment
dec $D$	$D \leftarrow D-1$	Decrement
neg $D$	$D \leftarrow -D$	Negate
not $D$	$D \leftarrow \sim D$	Complement
add $S, D$	$D \leftarrow D + S$	Add
sub $S, D$	$D \leftarrow D - S$	Subtract
imul $S, D$	$D \leftarrow D * S$	Multiply
xor $S, D$	$D \leftarrow D \sim S$	Exclusive-or
or $S, D$	$D \leftarrow D \mid S$	Or
and $S, D$	$D \leftarrow D \& S$	And
sall $k, D$	$D \leftarrow D \ll k$	Left shift
shl $k, D$	$D \leftarrow D \ll k$	Left shift (same as sall)
sar $k, D$	$D \leftarrow D \gg_A k$	Arithmetic right shift
shr $k, D$	$D \leftarrow D \gg_L k$	Logical right shift

CF	(unsigned) $t < (\text{unsigned}) a$	Unsigned overflow
ZF	$(t == 0)$	Zero
SF	$(t < 0)$	Negative
OF	$(a < 0 == b < 0) \&\& (t < 0 != a < 0)$	Signed overflow

CMP S1, S2 (S2 - S1)

TST S1, S2 (S1 & S2)

Instruction	Synonym	Effect	Set condition
sete $D$	setz	$D \leftarrow ZF$	Equal / zero
setne $D$	setnz	$D \leftarrow \sim ZF$	Not equal / not zero
sets $D$		$D \leftarrow SF$	Negative
setns $D$		$D \leftarrow \sim SF$	Nonnegative
setg $D$	setnle	$D \leftarrow \sim (SF \sim OF) \& \sim ZF$	Greater (signed >)
setge $D$	setnl	$D \leftarrow \sim (SF \sim OF)$	Greater or equal (signed >=)
setl $D$	setnge	$D \leftarrow SF \sim OF$	Less (signed <)
setle $D$	setng	$D \leftarrow (SF \sim OF) \mid ZF$	Less or equal (signed <=)
seta $D$	setnbe	$D \leftarrow \sim CF \& \sim ZF$	Above (unsigned >)
setae $D$	setnb	$D \leftarrow \sim CF$	Above or equal (unsigned >=)
setb $D$	setnae	$D \leftarrow CF$	Below (unsigned <)
setbe $D$	setna	$D \leftarrow CF \mid ZF$	Below or equal (unsigned <=)

jmp $Label$	1	Direct jump
jmp $*Operand$	1	Indirect jump

Callee Saved – Responsibility of function called to save the value before returning to function that called it. 1-6th arguments are Caller Saved

A struct within a struct gets aligned by largest value of struct

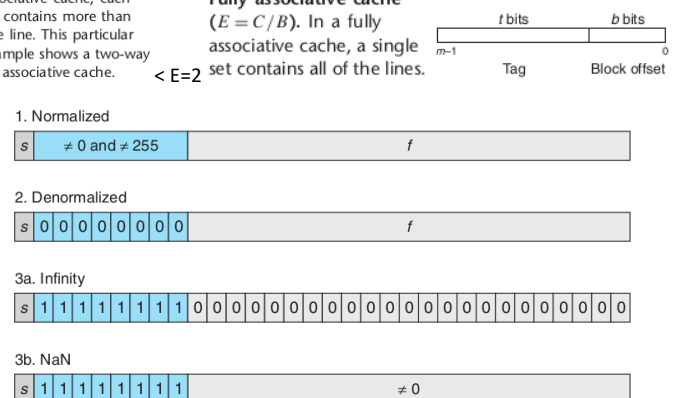
Parameter	Description
Fundamental parameters	
$S = 2^s$	Number of sets
$E$	Number of lines per set
$B = 2^b$	Block size (bytes)
$m = \log_2(M)$	Number of physical (main memory) address bits
Derived quantities	
$M = 2^m$	Maximum number of unique memory addresses
$s = \log_2(S)$	Number of <i>set index bits</i>
$b = \log_2(B)$	Number of <i>block offset bits</i>
$t = m - (s + b)$	Number of <i>tag bits</i>
$C = B \times E \times S$	Cache size (bytes), not including overhead such as

**Set associative cache**  
( $1 < E < C/B$ ). In a set associative cache, each set contains more than one line. This particular example shows a two-way set associative cache.

Figure 6.35

**Fully associative cache**  
( $E = C/B$ ). In a fully associative cache, a single set contains all of the lines.

The entire cache is one set, so by default set 0 is always selected.



For normalized numbers: For denormalized numbers:

$M = 1.xxxx$

$M = 0.xxxx$

$E = \text{exp} - \text{bias}$

$E = 1 - \text{bias}$

$\text{Bias} = 2^{(k-1)-1}$

$V = (-1)^s \cdot M \cdot 2^E$

## Floating Point: Rounding

1.BGRXXX

In the below examples, imagine the underlined part as a fraction.

- Guard Bit: the least significant bit of the resulting number
- Round Bit: the first bit removed from rounding
- Sticky Bits: all bits after the round bit, OR'd together

Examples of rounding cases, including rounding to nearest even number

- 1.10:11: More than  $\frac{1}{2}$ , round up: 1.11
- 1.10:10: Equal to  $\frac{1}{2}$ , round down to even: 1.10
- 1.01:01: Less than  $\frac{1}{2}$ , round down: 1.01
- 1.01:10: Equal to  $\frac{1}{2}$ , round up to even: 1.10
- 1.01:00: Equal to 0, do nothing: 1.01
- 1.00:00: Equal to 0, do nothing: 1.00

All other cases involve either rounding up or down - try them!

## Bonus Coverage: Arrays

Good toy examples (for your cheatsheet and/or big brain):



- A can be used as the pointer to the first array element: A[0]

	Type	Value
val	int *	x
val[2]	int	2
*(val + 2)	int	2
&val[2]	int *	x + 8
val + 2	int *	x + 8
val + i	int *	x + (4 * i)

Accessing methods:  
 • val[index]  
 • \*(val + index)

Addressing methods:  
 • &val[index]  
 • val + index


## Bonus Coverage: Arrays

Nested indexing rules (for your cheatsheet and/or big brain):

A[i][j] is element of type T, which requires K bytes


$$\text{Address } A + i * (C * K) + j * K$$

$$= A + (i * C + j) * K$$



## Bonus Coverage: Arrays

### Nested Array Element Access Code



```
leaq (%rdi,%rdi,4), %rax # 5*index
addl %rax,%rsi # 5*index+dig
movl pgh(,%rsi,4), %eax # M[pgh + 4*(5*index+dig)]
```

#### Array Elements

- pgh[index][dig] is int
- Address: pgh + 20\*index + 4\*dig  
= pgh + 4\*(5\*index + dig)

```
1 char c = 0xF0;
2 if(0xF0 == c){
3     //Never Enters unless GCCd
4     //c == 0xFFFFFFFF0;
5 }
1 printf("%x");
2 //Does not include 0x
3 printf("0x%.8x \n")
4 //Prints in Little Endian
5 printf("0x%x \n")
6 //Prints in Little Endian
```


1 Floats are four bytes!  
 2 Addresses on the stack are in order  
 3 All other data types on stack little endian (pointers too)  
 4 ! -> turns thing into signed ints  
 5 Size then signage

Struct  
 int;  
 float;  
 i=6942  
 f=0.4200  
 42, 69, 00, 00,  
 00, 42, 00, 00,

- 16 lines, 4-byte cache line size

- Physically addressed

- Direct mapped



$$T2U_w(x) = \begin{cases} x + 2^w, & x < 0 \\ x, & x \geq 0 \end{cases}$$

$$B2U_w(T2B_w(x)) = T2U_w(x) = x + x_{w-1}2^w$$

$$U2T_w(u) = \begin{cases} u, & u \leq TMax_w \\ u - 2^w, & u > TMax_w \end{cases}$$

$$U2T_w(u) = -u_{w-1}2^w + u$$

$$x *_w^u y = (x \cdot y) \bmod 2^w$$

$$x *_w^1 y = U2T_w((x \cdot y) \bmod 2^w)$$

PRINCIPLE: Unsigned division by a power of 2

For C variables x and k with unsigned values x and k, such that  $0 \leq k < w$ , the C expression  $x \gg k$  yields the value  $\lfloor x/2^k \rfloor$ .


PRINCIPLE: Two's-complement division by a power of 2, rounding up

Let C variables x and k have two's-complement value x and unsigned value k, respectively, such that  $0 \leq k < w$ . The C expression  $(x + (1 \ll k) - 1) \gg k$ , when the shift is performed arithmetically, yields the value  $\lceil x/2^k \rceil$ .

## Review of Symbols


### Basic Parameters

- N =  $2^n$ : Number of addresses in virtual address space
- M =  $2^m$ : Number of addresses in physical address space
- P =  $2^p$ : Page size (bytes)



### Components of the virtual address (VA)


- TLBI: TLB index
- TLBT: TLB tag
- VPO: Virtual page offset
- VPN: Virtual page number



### Components of the physical address (PA)

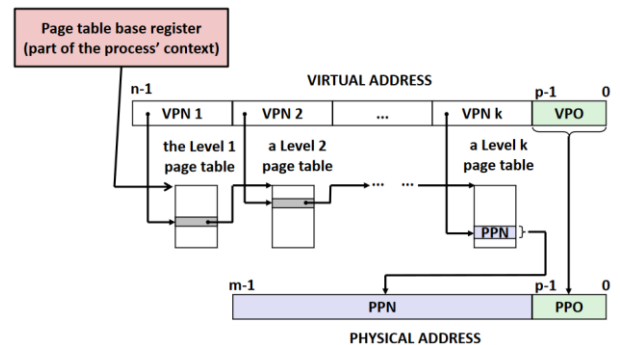
- PPO: Physical page offset (same as VPO)
- PPN: Physical page number
- CO: Byte offset within cache line
- CI: Cache index
- CT: Cache tag

(bits per field for our simple example)




## Translating with a k-level Page Table

- Having multiple levels greatly reduces page table size



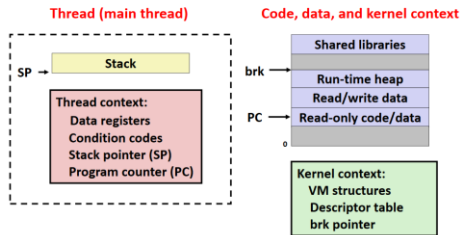
- 16 entries

- 4-way associative



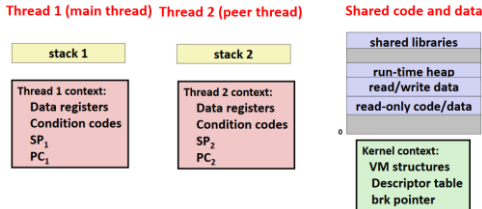
## Processes and such

- Process = thread + code, data, and kernel context




## A Process With Multiple Threads

- Multiple threads can be associated with a process
  - Each thread has its own logical control flow
  - Each thread shares the same code, data, and kernel context
  - Each thread has its own stack for local variables
    - but not protected from other threads
  - Each thread has its own thread id (TID)



- Separation of data is not strictly enforced:
  - Register values are truly separate and protected, but...
  - Any thread can read and write the stack of any other thread

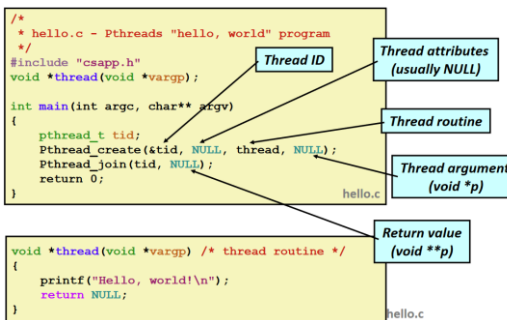


```
long* p = Malloc(sizeof(long));
*p = i;
Pthread_create(&tid[i],
              NULL,
              thread,
              (void *)p);

(i = 0; i < N; i++)
```

- Use malloc to create a per thread heap allocated place in memory for the argument

- Creating and reaping threads
  - pthread\_create()
  - pthread\_join()
- Determining your thread ID
  - pthread\_self()
- Terminating threads
  - pthread\_cancel()
  - pthread\_exit()
  - exit() [terminates all threads]
  - return [terminates current thread]
- Synchronizing access to shared variables
  - pthread\_mutex\_init
  - pthread\_mutex\_unlock



- Global variables
  - Def: Variable declared outside of a function
  - Virtual memory contains exactly one instance of any global variable
- Local variables
  - Def: Variable declared inside function without static attribute
  - Each thread stack contains one instance of each local variable
- Local static variables
  - Def: Variable declared inside function with the static attribute
  - Virtual memory contains exactly one instance of any local static variable.

## Shared Variable Analysis

- Which variables are shared?

Variable	Referenced by main thread?	Referenced by peer thread 0?	Referenced by peer thread 1?
ptr	yes	yes	yes
cnt	no	yes	yes
i, m	yes	no	no
msgs, m	yes	yes	yes
myid, p0	no	yes	no
myid, p1	no	no	yes

```
char **ptr; /* global var */
int main(int main, char *argv[]) {
    long i; pthread_t tid;
    char *msgs[2] = {"Hello from foo",
                    "Hello from bar"};

    ptr = msgs;
    for (i = 0; i < 2; i++)
        Pthread_create(&tid,
                       NULL, thread, (void *)i);
    Pthread_exit(NULL);
}

void *thread(void *vargp)
{
    long myid = (long)vargp;
    static int cnt = 0;
    printf("[%ld]: %s (cnt=%d)\n",
           myid, ptr[myid], ++cnt);
    return NULL;
}
```

- ptr, cnt, and msgs are shared
- i and myid are **not** shared

## Issues With Thread-Based Servers

- Must run "detached" to avoid memory leak
  - At any point in time, a thread is either joinable or detached
  - Joinable thread can be reaped and killed by other threads
    - must be reaped (with pthread\_join) to free memory resources
  - Detached thread cannot be reaped or killed by other threads
    - resources are automatically reaped on termination
  - Default state is joinable
    - use pthread\_detach(pthread\_self()) to make detached
- Must be careful to avoid unintended sharing
  - For example, passing pointer to main thread's stack
    - Pthread\_create(&tid, NULL, thread, (void \*)&connfd);
- All functions called by a thread must be **thread-safe**

## Appendix: Writing signal handlers

- G1. Call only async-signal-safe functions in your handlers.
  - Do not call printf, sprintf, malloc, exit! Doing so can cause deadlocks, since these functions may require global locks.
  - We've provided you with `sio_printf` which you can use instead.
- G2. Save and restore `errno` on entry and exit.
  - If not, the signal handler can corrupt code that tries to read `errno`.
  - The driver will print a warning if `errno` is corrupted.
- G3. Temporarily block signals to protect shared data.
  - This will prevent race conditions when writing to shared data.
- Avoid the use of global variables in `tslab`.

## Synchronization: Concurrency Issues

### Avoiding Deadlock Acquire shared resources in same order

```
int main(int argc, char** argv)
{
    pthread_t tid[2];
    Sem_init(&mutex[0], 0, 1); /* mutex[0] = 1 */
    Sem_init(&mutex[1], 0, 1); /* mutex[1] = 1 */
    Pthread_create(&tid[0], NULL, count, (void*) 0);
    Pthread_create(&tid[1], NULL, count, (void*) 1);
    Pthread_join(tid[0], NULL);
    Pthread_join(tid[1], NULL);
    printf("cnt=%d\n", cnt);
    return 0;
}

void *count(void *vargp)
{
    int i;
    int id = (int) vargp;
    for (i = 0; i < NITERS; i++) {
        P(&mutex[id]); P(&mutex[1-id]);
        cnt++;
        V(&mutex[id]); V(&mutex[1-id]);
    }
    return NULL;
}
```


```
Tid[0]:
P(s0);
P(s1);
cnt++;
V(s0);
V(s1);

Tid[1]:
P(s0);
P(s1);
cnt++;
V(s1);
V(s0);
```

- Classes of thread-unsafe functions:
  - Class 1: Functions that do not protect shared variables
  - Class 2: Functions that keep state across multiple invocations
  - Class 3: Functions that return a pointer to a static variable
  - Class 4: Functions that call thread-unsafe functions
- Failing to protect shared variables
  - Fix: Use P and V semaphore operations (or mutex)
  - Example: `goodcnt.c`
  - Issue: Synchronization operations will slow down code

## Reentrant Functions

- Def: A function is **reentrant** iff it accesses no shared variables when called by multiple threads.
  - Important subset of thread-safe functions
    - Require no synchronization operations
    - Only way to make a Class 2 function thread-safe is to make it reentrant (e.g., `rand_r`)



## Semaphores

- Semaphore: non-negative global integer synchronization variable

- Manipulated by P and V operations:
  - P(s): [ while (s == 0) wait(); s--; ]
    - Dutch for "Proberen" (test)
  - V(s): [ s++; ]
    - Dutch for "Verhogen" (increment)

- OS kernel guarantees that operations between brackets [ ] are executed indivisibly
  - Only one P or V operation at a time can modify s.
  - When while loop in P terminates, only that P can decrement s

- Semaphore invariant: (s >= 0)

- Mutex: binary semaphore used for mutual exclusion
  - P operation: "locking" the mutex
  - V operation: "unlocking" or "releasing" the mutex
  - "Holding" a mutex: locked and not yet unlocked.

- Define and initialize a mutex for the shared variable cnt:

```
volatile long cnt = 0; /* Counter */
sem_t mutex;          /* Semaphore that protects cnt */
sem_init(&mutex, 0, 1); /* mutex = 1 */
```

- Surround critical section with P and V:

```
for (i = 0; i < niters; i++) {
    P(&mutex);
    cnt++;
    V(&mutex);
}
```

```
linux> ./goodcnt 10000
OK cnt=20000
linux> ./goodcnt 10000
OK cnt=20000
linux>
```

- Define and initialize a mutex for the shared variable cnt:


```
volatile long cnt = 0; /* Counter */
pthread_mutex_t mutex;
pthread_mutex_init(&mutex, NULL); // No special attributes
```

- Surround critical section with lock and unlock:

```
for (i = 0; i < niters; i++) {
    pthread_mutex_lock(&mutex);
    cnt++;
    pthread_mutex_unlock(&mutex);
}
```

```
linux> ./goodcnt 10000
OK cnt=20000
linux> ./goodcnt 10000
OK cnt=20000
linux>
```

## Producer-Consumer Problem



- Common synchronization pattern:
  - Producer waits for empty slot, inserts item in buffer, and notifies consumer
  - Consumer waits for item, removes it from buffer, and notifies producer

## Circular Buffer (n = 10)

- Store elements in array of size n
- items: number of elements in buffer
- Empty buffer:
  - front = rear
- Nonempty buffer
  - rear: index of most recently inserted element
  - front: (index of next element to remove - 1) mod n
- Initially:

```
front 0
rear 0
items 0
```

0	1	2	3	4	5	6	7	8	9

```
void sbuf_init(sbuf_t *sp, int n)
{
    sp->buf = Calloc(n, sizeof(int));
    sp->n = n;
    sp->front = sp->rear = 0; /* Buffer holds max of n items */
    Sem_init(&sp->mutex, 0, 1); /* Empty buffer iff front == rear */
    Sem_init(&sp->slots, 0, n); /* Binary semaphore for locking */
    Sem_init(&sp->items, 0, 0); /* Initially, buf has n empty slots */
    Sem_init(&sp->items, 0, 0); /* Initially, buf has zero items */
}
```

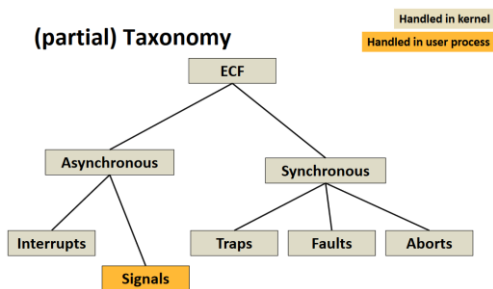
```
/* Insert item onto the rear of shared buffer sp */
void sbuf_insert(sbuf_t *sp, int item)
{
    P(&sp->slots); /* Wait for available slot */
    P(&sp->mutex); /* Lock the buffer */
    if (++sp->rear >= sp->n) /* Increment index (mod n) */
        sp->rear = 0;
    sp->buf[sp->rear] = item; /* Insert the item */
    V(&sp->mutex); /* Unlock the buffer */
    V(&sp->items); /* Announce available item */
}
```

```
/* Remove and return the first item from buffer sp */
int sbuf_remove(sbuf_t *sp)
{
    int item;
    P(&sp->items); /* Wait for available item */
    P(&sp->mutex); /* Lock the buffer */
    if (++sp->front >= sp->n) /* Increment index (mod n) */
        sp->front = 0;
    item = sp->buf[sp->front]; /* Remove the item */
    V(&sp->mutex); /* Unlock the buffer */
    V(&sp->slots); /* Announce available slot */
    return item;
}
```




## ECF

### (partial) Taxonomy




- An **exception** is a transfer of control to the OS *kernel* in response to some **event** (i.e., change in processor state)
  - Kernel is the memory-resident part of the OS
  - Examples of events: Divide by 0, arithmetic overflow, page fault, I/O request completes, typing Ctrl-C



- User calls: `open(filename, options)`
- Calls `_open` function, which invokes system call `syscall`


```

0000000000e5d70 <_open>:
...
e5d79: b8 02 00 00 00 mov $0x2,%eax # open is syscall #2
e5d7e: 0f 05 syscall # Return value in %rax
e5d80: 48 3d 01 10 ff cmp $0xfffffffff001,%rax
...
e5dfa: c3 retq
    
```



- `%rax` contains syscall number
- Other arguments in `%rdi`, `%rsi`, `%rdx`, `%r10`, `%r8`, `%r9`
- Return value in `%rax`
- Negative value is an error corresponding to negative `errno`

## Conceptual View of fork



## WAITPID

```

pid_t wait(int *stat_loc);
pid_t waitpid(pid_t pid, int *stat_loc, int options);
    
```

If `pid` is equal to `(pid_t)-1`, `status` is requested for any child process. In this respect, `waitpid()` is then equivalent to `wait()`.

If `pid` is greater than 0, it specifies the process ID of a single child process for which `status` is requested.

If `pid` is 0, `status` is requested for any child process whose process group ID is equal to that of the calling process.

If `pid` is less than `(pid_t)-1`, `status` is requested for any child process whose process group ID is equal to the absolute value of `pid`.

`WIFEXITED(stat_val)`

Evaluates to a non-zero value if `status` was returned for a child process that terminated normally.

`WEXITSTATUS(stat_val)`

If the value of `WIFEXITED(stat_val)` is non-zero, this macro evaluates to the low-order 8 bits of the `status` argument that the child process passed to `_exit()` or `exit()`, or the value the child process returned from `main()`.

`WIFSIGNALED(stat_val)`

Evaluates to a non-zero value if `status` was returned for a child process that terminated due to the receipt of a signal that was not caught (see <signal.h>).

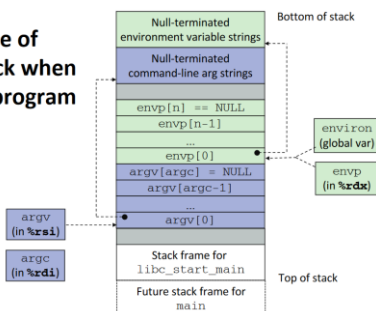
## Return Value

If `wait()` or `waitpid()` returns because the status of a child process is available, these functions shall return a value equal to the process ID of the child process for which `status` is reported. If `wait()` or `waitpid()` returns due to the delivery of a signal to the calling process, -1 shall be returned and `errno` set to `[EINTR]`. If `waitpid()` was invoked with `WNOHANG` set in `options`, it has at least one child process specified by `pid` for which `status` is not available, and `status` is not available for any process specified by `pid`, 0 is returned. Otherwise, `(pid_t)-1` shall be returned, and `errno` set to indicate the error.

## EXECVE

- `int execve(char *filename, char *argv[], char *envp[])`
- Loads and runs in the current process:
  - Executable file `filename`
    - Can be object file or script file beginning with `#!/interpreter` (e.g., `#!/bin/bash`)
  - ...with argument list `argv`
    - By convention `argv[0] == filename`
  - ...and environment variable list `envp`
    - "name=value" strings (e.g., `USER=droh`)
    - `getenv`, `putenv`, `printenv`
- Overwrites code, data, and stack
  - Retains PID, open files and signal context
- Called **once** and **never** returns
  - ...except if there is an error

## Structure of the stack when a new program starts



## Signal Handling

- The `signal` function modifies the default action associated with the receipt of signal `signal`:
  - `handler_t *signal(int signal, handler_t *handler)`
- Different values for `handler`:
  - `SIG_IGN`: ignore signals of type `signal`
  - `SIG_DFL`: revert to the default action on receipt of signals of type `signal`
  - Otherwise, `handler` is the address of a user-level **signal handler**
- Implicit blocking mechanism
  - Kernel blocks any pending signals of type currently being handled
  - e.g., a `SIGINT` handler can't be interrupted by another `SIGINT`
- Explicit blocking and unblocking mechanism
  - `sigprocmask` function

```

sigset_t mask, prev_mask;

Sigemptyset(&mask);
Sigaddset(&mask, SIGINT);

/* Block SIGINT and save previous blocked set */
Sigprocmask(SIG_BLOCK, &mask, &prev_mask);

/* Code region that will not be interrupted by SIGINT */

/* Restore previous blocked set, unblocking SIGINT */
Sigprocmask(SIG_SETMASK, &prev_mask, NULL);
    
```

## Guidelines for Writing Safe Handlers

- G0: Keep your handlers as simple as possible
  - e.g., set a global flag and return
- G1: Call only **async-signal-safe** functions in your handlers
  - `printf`, `fprintf`, `malloc`, and `exit` are not safe!
- G2: Save and restore `errno` on entry and exit
  - So that other handlers don't overwrite your value of `errno`
- G3: Protect accesses to shared data structures by temporarily blocking all signals
  - To prevent possible corruption
- G4: Declare global variables as **volatile**
  - To prevent compiler from storing them in a register
- G5: Declare global flags as **volatile sig\_atomic\_t**
  - `flag`: variable that is only read or written (e.g. `flag = 1`, not `flag++`)
  - Flag declared this way does not need to be protected like other globals
- Function is **async-signal-safe** if either reentrant (e.g., all variables stored on stack frame, CS:APP3e 12.7.2) or non-interruptible by signals
- Posix guarantees 117 functions to be **async-signal-safe**
  - Source: "man 7 signal-safety"
  - Popular functions on the list:
    - `_exit`, `write`, `wait`, `waitpid`, `sleep`, `kill`

```

while (n--) {
    Sigprocmask(SIG_BLOCK, &mask_one, &prev_one); /* Block SIGCHLD */
    if ((pid = Fork()) == 0) { /* Child process */
        Sigprocmask(SIG_SETMASK, &prev_one, NULL); /* Unblock SIGCHLD */
        Execve("/bin/date", argv, NULL);
    }
    Sigprocmask(SIG_BLOCK, &mask_all, NULL); /* Parent process */
    addjob(pid); /* Add the child to the job list */
    Sigprocmask(SIG_SETMASK, &prev_one, NULL); /* Unblock SIGCHLD */
}
exit(0);
    
```

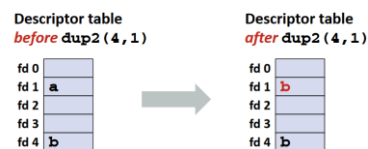
- `int sigsuspend(const sigset_t *mask)`
- Equivalent to atomic (uninterruptable) version of:

```

sigprocmask(SIG_SETMASK, &mask, &prev);
pause();
sigprocmask(SIG_SETMASK, &prev, NULL);
    
```

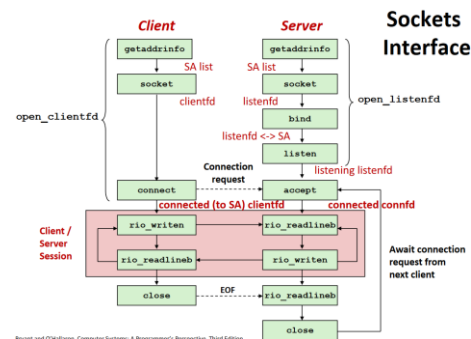
## FILE IO

- Answer: By calling the `dup2(oldfd, newfd)` function
  - Copies (per-process) descriptor table entry `oldfd` to entry `newfd`




- When to use standard I/O
  - When working with disk or terminal files
- When to use raw Unix I/O
  - Inside signal handlers, because Unix I/O is **async-signal-safe**

## Network



## Review: getaddrinfo Linked List



- Clients: walk this list, trying each socket address in turn, until the calls to `socket` and `connect` succeed.
- Servers: walk the list until calls to `socket` and `bind` succeed.

## OTHER

**Snoopy Caches**

- Invalid: cannot use value
- Shared: readable copy
- Exclusive: writable copy

start with E → S

**VM locality**

- working set size < main memory
- good performance
- working set size > main memory
- thrashing

KB → 2<sup>10</sup>

- deadlock: nothing can proceed (need locks)
- livelock: infinite loop
- starvation: unfairness, one has priority

## Garbage collection

- mark & sweep: mark all connected to roots
- sweep all allocated and not marked.